

# Homework 11

Fangda Li  
li1208@purdue.edu  
ECE 661 - Computer Vision

January 5, 2017

## 1 Methodology

This section briefly describes the principle steps of PCA, LDA and cascaded AdaBoost for classification problems.

### 1.1 PCA (Principle Component Analysis)

PCA aims to find the optimal subspace (of the original feature space) onto which the projected samples have the highest variance. The bases of the subspace are the first  $K$  eigenvectors of the covariance matrix  $\mathbf{C}$ , which is defined by

$$\mathbf{C} = \frac{1}{N} \sum_{i=0}^{N-1} (\vec{x}_i - \vec{m})(\vec{x}_i - \vec{m})^T. \quad (1)$$

where  $\vec{x}_i$  The basis of the subspace is calculated as follows:

1. Suppose  $\mathbf{X} \in R_{M,N,(M>N)}$ , represents the  $N$  training samples each with  $M$  features. Moreover,  $\mathbf{X}$  is normalized such that it has zero mean and unit variance.
2. The covariance matrix can be rewritten as

$$\mathbf{C} = \mathbf{X}\mathbf{X}^T. \quad (2)$$

3. Compute the eigenvectors  $\vec{u}$  for  $\mathbf{X}^T\mathbf{X}$ . Then the eigenvectors  $\vec{w}$  for  $\mathbf{X}\mathbf{X}^T$  are obtained by

$$\vec{w} = \mathbf{X}\vec{u}. \quad (3)$$

4. Since the rank of  $\mathbf{X}\mathbf{X}^T$  is  $N$ , at most  $N$  eigenvectors  $\vec{w}$  can be obtained. The final basis of the subspace is choose to be the  $K$  eigenvectors corresponding to the  $K$  largest eigenvalues, where  $K \leq N$ .

After the basis has been obtained, all training samples are projected to the optimal subspace. Next, nearest neighbor approach is used in the optimal subspace to determine the label of a test sample.

## 1.2 LDA (Linear Discriminant Analysis)

LDA aims to find the most discriminating subspace in terms of the ratio of between-class scatter and within-class scatter. The steps of LDA for a multi-class discriminant problem are described as following:

1. Define the within-class scatter:

$$\mathbf{S}_W = \frac{1}{C} \sum_{i=1}^C \frac{1}{C_i} \sum_{k=1}^{C_i} (\vec{x}_k^i - \vec{m}_i)(\vec{x}_k^i - \vec{m}_i)^T, \quad (4)$$

where  $x_k^i$  is the  $k$ th sample of the  $i$ th class.

2. Define the between-class scatter:

$$\mathbf{S}_B = \frac{1}{C} \sum_{i=1}^C (\vec{m}_i - \vec{m})(\vec{m}_i - \vec{m})^T. \quad (5)$$

3. The most discriminating vector will therefore maximize the Fisher Discriminant Function:

$$J(\vec{w}) = \frac{\vec{w}^T \mathbf{S}_B \vec{w}}{\vec{w}^T \mathbf{S}_W \vec{w}}. \quad (6)$$

4. Finally, the calculation of the LDA eigenvectors is carried out by calculating the  $C - 1$  mean difference vectors  $\vec{m}_i - \vec{m}$ .

Similar to PCA, the final classification task is done with the nearest neighbor approach on the projections onto the optimal subspace.

## 1.3 Cascaded AdaBoost

The overall idea of Viola and Jones is to cascade independently trained AdaBoost classifiers to suppress the false positive rate while preserving the recall. Several key insights of the algorithm are listed below:

1. Each cascade of the Viola and Jones classifier is an AdaBoost classifier. The main intuition of AdaBoost is that, after selecting the best weak classifier within each iteration, the weights of the misclassified samples by the new weak classifier are boosted. As a result, in the next iteration, the weak classifiers are evaluated with previously misclassified samples having larger weights. The concept is similar to hard negative mining in a sense. Moreover, AdaBoost assigns confidence values to each weak classifier based on its own error. Prediction of the final strong classifier is a weighted vote from all the previous best weak classifiers.
2. For each AdaBoost classifier in the cascade, it is bound to achieve a certain detection rate  $d$  while keeping its false positive rate  $f$  below a certain value. Then, the overall cascade classifier will have a false positive rate  $F$ ,

$$F = \prod_{i=1}^K f_i, \quad (7)$$

and a detection rate  $D$ ,

$$D = \prod_{i=1}^K d_i. \quad (8)$$

As a result of the products, the rate of decrease in  $F$  is much slower than  $D$ , given  $D$  is close to 1 and  $F$  is close to 0.

3. The full training procedure for one strong AdaBoost classifier is shown as followed:
  - (a) Given  $m$  positive samples ( $y = 1$ ) and  $l$  negative samples ( $y = 0$ ), initialize the weights,  $w_{1,i} = \frac{1}{2m}, \frac{1}{2l}$ , for every positive and negative sample, respectively.
  - (b) For  $t = 1, \dots, T$ :
    - i. Normalize weights of all samples,  $w_{t,i} = \frac{w_{t,i}}{\sum_j w_{t,i}}$ .
    - ii. Selected best weak classifier with respect to the weighted error

$$\epsilon_t = \min_{f,p,\theta} \sum_i w_i |h(x_i, f, p, \theta) - y_i|. \quad (9)$$

- iii. Construct new weak classifier  $h_t(x) = h(x, f, p, \theta)$  with  $f, p, \theta$  minimizing  $\epsilon_t$ .
  - iv. Update weights for the samples using  $w_{t+1,i} = w_{t,i} \beta_t^{1-e_i}$ , where  $e_i = 1$  when  $x_i$  is classified incorrectly and 0 otherwise,  $\beta_t = \frac{\epsilon_t}{1-\epsilon_t}$  is the confidence of the current weak classifier.
- (c) The final strong AdaBoost classifier is:

$$C(x) = \begin{cases} 1 & \sum_t \alpha_t h_t(x) \geq 0.5 \sum_t \alpha_t \\ 0 & \text{otherwise} \end{cases} \quad (10)$$

where  $\alpha_t = \log \frac{1}{\beta_t}$ .

## 2 Results and Discussion

### 2.1 PCA and LDA

PCA and LDA are evaluated simultaneously on the face dataset. The performance can be found in Figure 1, in which LDA slightly outperforms PCA given the same dimension of subspace. Nevertheless, both PCA and LDA are able to significantly reduce the dimensionality of the feature space while staying discriminating.




Figure 1: Classification accuracy v.s. Dimension of subspace

## 2.2 Cascaded AdaBoost

This subsection show the results of the Viola and Jones framework on the car dataset. The dataset consists of 710 p




Figure 2: Detection Rate v.s. Number of Stages




Figure 3: False Positive Rate v.s. Number of Stages




Figure 4: Accuracy v.s. Number of Stages

## 3 Source Code

### 3.1 pca.py

```
#!/usr/bin/python
from pylab import *
import cv2
from sklearn.neighbors import NearestNeighbors

class PCAClassifier(object):
    def __init__(self):
        super(PCAClassifier, self).__init__()
        self.K = 0
        self.m = None
        self.WK = None
        self.NN = None

    def train(self, train_data, train_label, K):
        """
        Construct K-D tree based on the projections onto the subspace
        """
        print "=====PCA Training====="
        self.train_data = train_data
        self.train_label = train_label
        self.K = K
```

```

        # Follow the notation of Avi's tutorial
        X = self.train_data
        m = mean(X, axis=1)
        _, _, Ut = svd(dot(X.T, X))
        W = dot(X, Ut.T)
        # Preserve the first K eigenvectors
        WK = W[:, :self.K]
        # Project all training samples to the K-D subspace
        Y = dot( WK.T, X - m.reshape(-1,1) )
        # Construct K-D tree
        self.NN = NearestNeighbors(n_neighbors=1, metric='euclidean').fit(Y.T)
        self.m = m
        self.WK = WK

def test(self, test_data, test_label):
    '''
        Project onto the K-D optimal subspace and find the nearest nei
    '''
    print "=====PCA Testing====="
    num_test = test_label.size
    # Projection first
    X = test_data
    Y = dot( self.WK.T, X - self.m.reshape(-1,1) )
    _, indices = self.NN.kneighbors(Y.T)
    pred = self.train_label[indices].flatten()
    accuracy = sum((pred - test_label) == 0) / float(num_test)
    print "K = {0}; Accuracy = {1:.2f}%".format(self.K, accuracy*100)
    return accuracy

```

## 3.2 lda.py

```

#!/usr/bin/python
from pylab import *
import cv2
from sklearn.neighbors import NearestNeighbors

class LDAClassifier(object):
    def __init__(self):
        super(LDAClassifier, self).__init__()
        self.K = 0
        self.W = None
        self.NN = None

    def train(self, train_data, train_label, K):
        '''

```

```

        Construct K-D tree based on the projections onto the subspace
    """
    print "===== LDA Training ====="
    self.train_data = train_data
    self.train_label = train_label
    self.K = K
    self.num_classes = unique(train_label).size
    # Follow Avi's notation
    X = self.train_data
    L = self.train_label
    C = self.num_classes
    if self.K > C-1: self.K = C-1
    # Get the class means and global mean
    m = mean(X, axis=1)
    M = zeros((X.shape[0],C))
    for i in range(1,C+1):
        M[:,i-1] = mean(X[:, L==i], axis=1)
    # Eigenvectors of the between class scatter are the same as mean difference
    W = M-m.reshape(-1,1)
    W = W[:, :self.K]
    # Project all training samples to the K-D subspace
    Y = dot( W.T, X - m.reshape(-1,1) )
    # Construct K-D tree
    self.NN = NearestNeighbors(n_neighbors=1, metric='euclidean').fit(Y.T)
    self.m = m
    self.W = W

def test(self, test_data, test_label):
    """
        Project onto the K-D optimal subspace and find the nearest neighbors
    """
    print "===== LDA Testing ====="
    num_test = test_label.size
    # Projection first
    X = test_data
    Y = dot( self.W.T, X - self.m.reshape(-1,1) )
    _, indices = self.NN.kneighbors(Y.T)
    pred = self.train_label[indices].flatten()
    accuracy = sum((pred - test_label) == 0) / float(num_test)
    print "Accuracy = {0:.2f}%".format(accuracy*100)
    return accuracy

```



### 3.3 adaboost.py

```
#!/usr/bin/python
from pylab import *
import cv2

WIDTH = 40
HEIGHT = 20
NUMFEATURES = 47232 # Obtained by _get_feature_matrix()

def get_integral_image(image):
    """
    Compute and return the integral representation of an gray-scale image.
    """
    return cumsum(cumsum(image,axis=0),axis=1)

def _get_feature_matrix():
    """
    Return the feature matrix.
    """
    features = zeros((NUMFEATURES, WIDTH*HEIGHT), dtype=int8)
    print "Constructing feature matrix...", features.shape
    count = 0
    # Offset from image boundary
    offset = 2
    # Do horizontal ones first, 1x2 base size
    bW, bH = 2, 1
    for i in range(1, HEIGHT, bH):
        for j in range(1, WIDTH, bW):
            for y in range(0 + offset, HEIGHT - bH*i + 1 - offset):
                for x in range(0 + offset, WIDTH - bW*j + 1 - offset):
                    features[count, y*HEIGHT+x] = 1.0
                    features[count, y*HEIGHT+x+bW*j/2] = -2.0
                    features[count, y*HEIGHT+x+bW*j] = 1.0
                    features[count, (y+bH*i)*HEIGHT+x] = -1.0
                    features[count, (y+bH*i)*HEIGHT+x+bW*j/2] = 2.0
                    features[count, (y+bH*i)*HEIGHT+x+bW*j] = -1.0
                    count = count + 1
    # Vertical features, 2x1 base size
    bW, bH = 1, 2
    for i in range(1, HEIGHT, bH):
        for j in range(1, WIDTH, bW):
            for y in range(0 + offset, HEIGHT - bH*i + 1 - offset):
                for x in range(0 + offset, WIDTH - bW*j + 1 - offset):
                    features[count, y*HEIGHT+x] = -1.0
```

```

        features[count, y*HEIGHT+x+bW*j] = 1.0
        features[count, (y+bH*i/2)*HEIGHT+x] = 2.0
        features[count, (y+bH*i/2)*HEIGHT+x+bW*j] = -2.0
        features[count, (y+bH*i)*HEIGHT+x/2] = -1.0
        features[count, (y+bH*i)*HEIGHT+x+bW*j] = 1.0
        count = count + 1

    print "Total number of features...", count
    return features

def _extract_features(data, feature_matrix):
    '''
        Extract all features from integral images into a nFeatures x nSamples
        Also return the sorted indices along the samples direction.
    '''
    # We are only concerned with two-rectangle edge features
    # Get the sorted list of sample indices.
    feature_vectors = dot( feature_matrix, data )
    print "Feature vectors size...", feature_vectors.shape
    print "Sorting samples based on feature values..."
    sorted_indices = argsort(feature_vectors, axis=1)
    print "Sorted indices size...", sorted_indices.shape
    return sorted_indices, feature_vectors

class CascadedAdaBoostClassifier(object):
    def __init__(self):
        super(CascadedAdaBoostClassifier, self).__init__()
        self.cascaded_adaboost = []
        self.train_data = None
        self.train_label = None
        self.train_feat_vecs = None
        self.train_num_pos = 0
        self.train_num_neg = 0
        self.test_data = None
        self.test_label = None
        self.feature_matrix = _get_feature_matrix()

    def set_testing_data(self, test_data, test_label):
        self.test_data = test_data
        self.test_label = test_label

    def train(self, train_data, train_label, num_stages, num_feats):
        # def train(self, train_data, train_label, f, d, Ftarg, maxIter):
        '''
            Train cascaded AdaBoost classifiers given user-defined:
            max acceptable fpr per layer f, min acceptable detection rate

```

```

        and overall fpr  $F_{target}$ .
    """
    self.train_feat_vecs = dot( self.feature_matrix, train_data )
    self.train_num_pos = int(sum(train_label))
    self.train_num_neg = train_label.size - self.train_num_pos
    self.train_num = train_label.size
    self.train_data, self.train_label = train_data, train_label
    # Positive and negative training samples for current stage
    all_pos_feat_vecs = self.train_feat_vecs[:,self.train_label==1]
    all_neg_feat_vecs = self.train_feat_vecs[:,self.train_label==0]
    pos_feat_vecs = all_pos_feat_vecs
    neg_feat_vecs = all_neg_feat_vecs
    Flog_train = []
    Dlog_train = []
    Alog_train = []
    Flog_test = []
    Dlog_test = []
    Alog_test = []
    # Add stages
    for i in range(num_stages):
        print "Training %dth AdaBoost classifier in the cascade..." % (i+1)
        current_adaboost = self._add_adaboost_classifier()
        current_adaboost.set_training_feature_vectors(pos_feat_vecs, neg_feat_vecs)
        # Add features
        for j in range(num_feats):
            print "Adding feature %d..." % (j+1)
            current_adaboost.add_weak_classifier()
        # Update negative samples to use
        fp_indices, F, D, A = self._classify_training_data()
        # Record training info
        Flog_train.append(F)
        Dlog_train.append(D)
        Alog_train.append(A)
        neg_feat_vecs = all_neg_feat_vecs[:,fp_indices-self.train_num_pos]
        # Record testing info
        F, D, A = self.test()
        Flog_test.append(F)
        Dlog_test.append(D)
        Alog_test.append(A)
        print "Training:"
        print "FP:\n", Flog_train
        print "RC:\n", Dlog_train
        print "AC:\n", Alog_train
        print "Testing:"
        print "FP:\n", Flog_test

```

```

        print "RC:\n", Dlog_test
        print "AC:\n", Alog_test

def _add_adaboost_classifier(self):
    '''
        Allocate and return a new AdaBoost classifier.
    '''
    c = AdaBoostClassifier()
    c.set_feature_matrix(self.feature_matrix)
    self.cascaded_adaboost.append(c)
    return c

def _classify_training_data(self):
    '''
        Evaluate the cascaded classifier on the training data,
        and return the indices of false positive samples as well as the
    '''
    print "Classifying training images..."
    feat_vecs = self.train_feat_vecs
    pos_indices = arange(self.train_num)
    for classifier in self.cascaded_adaboost:
        preds = classifier.classify_feature_vectors(feat_vecs)
        # Only pass on the samples with positive predictions
        feat_vecs = feat_vecs[:,preds==1]
        pos_indices = pos_indices[preds==1]
    # Final prediction
    fp_indices = pos_indices[ self.train_label[pos_indices] == 0 ]
    num_tp = sum(self.train_label[pos_indices])
    D = num_tp*1.0 / self.train_num_pos
    # Calculate false positive rate by counting the zeros
    F = (pos_indices.size - num_tp)*1.0 / self.train_num_neg
    w = self.train_num_pos*1.0 / (self.train_num_pos + self.train_num_neg)
    A = D * w + (1-F)*(1-w)
    print "F = %.4f, D = %.4f, A = %.4f" % (F,D,A)
    return fp_indices, F, D, A

def test(self):
    '''
        Classify test images.
    '''
    print "Classifying testing images..."
    feat_vecs = dot( self.feature_matrix, self.test_data )
    test_num_pos = int(sum(self.test_label))
    test_num_neg = self.test_label.size - test_num_pos
    test_num = self.test_label.size

```

```

pos_indices = arange(test_num)
for classifier in self.cascaded_adaboost:
    preds = classifier.classify_feature_vectors(feet_vecs)
    # Only classify the samples with postive predictions
    feat_vecs = feat_vecs[:,preds==1]
    pos_indices = pos_indices[preds==1]
# Final prediction
num_tp = sum(self.test_label[pos_indices])
D = num_tp*1.0 / test_num_pos
# Calculate false positive rate by counting the zeros
F = (pos_indices.size - num_tp)*1.0 / test_num_neg
w = test_num_pos*1.0 / (test_num_pos + test_num_neg)
A = D * w + (1-F)*(1-w)
print "F = %.4f, D = %.4f, A = %.4f" % (F,D,A)
return F,D,A

class AdaBoostClassifier(object):
    def __init__(self):
        super(AdaBoostClassifier, self).__init__()
        self.train_label = None
        self.train_sorted_indices = None
        self.train_feat_vecs = None
        self.train_num_pos = 0
        self.train_num_neg = 0
        self.threshold = 1.0
        self.sample_weights = None
        self.weak_classifier_indices = array([], dtype=int)
        self.weak_classifier_polarities = array([])
        self.weak_classifier_threshs = array([])
        self.weak_classifier_weights = array([])
        self.weak_classifier_results = array([])
        self.weak_classifier_weighted_results = None

    def set_feature_matrix(self, feature_matrix):
        self.feature_matrix = feature_matrix

    def set_training_feature_vectors(self, pos_feat_vecs, neg_feat_vecs):
        '''
                Given current training feature vectors, sort them.
        '''

        self.train_num_pos = pos_feat_vecs.shape[1]
        self.train_num_neg = neg_feat_vecs.shape[1]
        self.train_label = hstack( (ones(self.train_num_pos), zeros(self.train_num_neg)) )
        self.train_feat_vecs = hstack( (pos_feat_vecs, neg_feat_vecs) )
        self.train_sorted_indices = argsort(self.train_feat_vecs, axis=1)

```

```

        print "Number of positive / negative samples in training...", self.train_label

def add_weak_classifier(self):
    """
        Add the current best weak classifier on the weighted training data
    """
    # Initialize all the weights if this is the first weak classifier
    if self.weak_classifier_indices.size == 0:
        self.sample_weights = zeros(self.train_label.size, dtype=float)
        self.sample_weights.fill( 1.0 / (2 * self.train_num_neg) )
        self.sample_weights[self.train_label==1] = 1.0 / (2 * self.train_num_pos)
    # Normalize the weights otherwise
    else:
        self.sample_weights = self.sample_weights / sum(self.sample_weights)
    # Now pick the weak classifier with the min error with respect to the training data
    best_feat_index, best_feat_polarity, best_feat_thresh, best_feat_error, best_feat_results = self._find_best_weak_classifier()
    # Update our list of weak classifiers
    self.weak_classifier_indices = append(self.weak_classifier_indices, best_feat_index)
    self.weak_classifier_polarities = append(self.weak_classifier_polarities, best_feat_polarity)
    self.weak_classifier_threshs = append(self.weak_classifier_threshs, best_feat_thresh)
    # Get confidence value of the best new classifier
    # Following the notation in the paper
    beta = best_feat_error / (1 - best_feat_error)
    alpha = log(1 / abs(beta))
    self.weak_classifier_weights = append(self.weak_classifier_weights, alpha)
    e = abs(best_feat_results - self.train_label)
    self.sample_weights = self.sample_weights * beta**(1-e)
    # Adjust the threshold
    if self.weak_classifier_results.size == 0:
        self.weak_classifier_results = best_feat_results.reshape(-1,1)
    else:
        self.weak_classifier_results = hstack((self.weak_classifier_results, best_feat_results.reshape(-1,1)))
    self.weak_classifier_weighted_results = dot(self.weak_classifier_results, self.weak_classifier_weights)
    self.threshold = min(self.weak_classifier_weighted_results[self.train_label==1])

def _get_best_weak_classifier(self):
    """
        Return the index of the best feature with the minimum weighted error
    """
    feature_errors = zeros(NUMFEATURES)
    feature_thresh = zeros(NUMFEATURES)
    feature_polarity = zeros(NUMFEATURES)
    feature_sorted_index = zeros(NUMFEATURES, dtype=int)
    Tplus = sum(self.sample_weights[self.train_label==1])
    Tminus = sum(self.sample_weights[self.train_label==0])

```

```

# Iterate to find the best feature
for r in range(NUMFEATURES):
    sorted_weights = self.sample_weights[self.train_sorted_indices[r]]
    sorted_labels = self.train_label[self.train_sorted_indices[r,:]]
    Splus = cumsum(sorted_labels * sorted_weights)
    Sminus = cumsum(sorted_weights) - Splus
    # Error of choice influences the polarity
    Eplus = Splus + Tminus - Sminus
    Eminus = Sminus + Tplus - Splus
    polarities = zeros(self.train_num_pos + self.train_num_neg)
    polarities[Eplus > Eminus] = -1
    polarities[Eplus <= Eminus] = 1
    errors = minimum(Eplus, Eminus)
    sorted_index = argmin(errors)
    min_error_sample_index = self.train_sorted_indices[r,sorted_index]
    min_error = min(errors)
    threshold = self.train_feat_vecs[r, min_error_sample_index]
    polarity = polarities[sorted_index]
    feature_errors[r] = min_error
    feature_thresh[r] = threshold
    feature_polarity[r] = polarity
    feature_sorted_index[r] = sorted_index

# Now pick the best one
best_feat_index = argmin(feature_errors)
best_feat_thresh = feature_thresh[best_feat_index]
best_feat_error = feature_errors[best_feat_index]
best_feat_polarity = feature_polarity[best_feat_index]
best_feat_results = zeros(self.train_num_pos + self.train_num_neg)
best_sorted_index = feature_sorted_index[best_feat_index]
if best_feat_polarity == 1:
    best_feat_results[ self.train_sorted_indices[best_feat_index, be
else:
    best_feat_results[ self.train_sorted_indices[best_feat_index, :b
print 'index, polarity, thresh, error'
print best_feat_index, best_feat_polarity, best_feat_thresh, best_feat_e
return best_feat_index, best_feat_polarity, best_feat_thresh, best_feat_e

def classify_feature_vectors(self, feat_vecs):
    '''
        Classify feature vectors and return classified labels.
    '''
    # Get the feature values
    weak_classifiers = feat_vecs[self.weak_classifier_indices,:]
    # Organize as column vectors to ease broadcasting later
    polar_colvec = self.weak_classifier_polarities.reshape(-1,1)

```

```

thresh_colvec = self.weak_classifier_threshs.reshape(-1,1)
# Predictions of all weak classifiers
weak_classifier_preds = weak_classifiers * polar_colvec > thresh_colvec
weak_classifier_preds[weak_classifier_preds==True] = 1
weak_classifier_preds[weak_classifier_preds==False] = 0
# Apply weak classifier weights
strong_classifier_result = dot(self.weak_classifier_weights, weak_classifi
# Apply strong classifier threshold
final_preds = zeros(strong_classifier_result.size)
final_preds[strong_classifier_result >= self.threshold] = 1
return final_preds

```

### 3.4 main.py

```

#!/usr/bin/python
from pylab import *
import cv2
import os
from pca import PCAClassifier
from lda import LDAClassifier
from adaboost import *

def load_face_dataset():
    '''
        Load the face dataset in the following format:
        - each image is converted to gray scale
        - each face image is vectorized as a column vector
        - labels are organized as a column vector
    '''
    print "Loading face dataset..."
    # Process training images first
    train_path = './face-dataset/train/'
    train_files = [f for f in os.listdir(train_path) if f.endswith(".png")]
    num_train = len(train_files)
    train_data = zeros((128*128, num_train), dtype=float)
    train_label = zeros(num_train, dtype=int)
    for i,f in enumerate(train_files):
        image = imread(os.path.join(train_path,f))
        image = cv2.cvtColor(image, cv2.COLOR_RGBA2GRAY)
        train_data[:,i] = image.flatten()
        train_label[i] = int(f.split('_')[0])
    # Normalization across all images, zero mean and unit variance
    train_mean = mean(train_data)
    train_std = std(train_data)
    train_data = train_data - train_mean

```



```

train_data = train_data / train_std
# Process testing images
test_path = './face-dataset/test/'
test_files = [f for f in os.listdir(test_path) if f.endswith(".png")]
num_test = len(test_files)
test_data = zeros((128*128, num_test), dtype=float)
test_label = zeros(num_test, dtype=int)
for i,f in enumerate(test_files):
    image = imread(os.path.join(test_path,f))
    image = cv2.cvtColor(image, cv2.COLOR_RGBA2GRAY)
    test_data[:,i] = image.flatten()
    test_label[i] = int(f.split('_')[0])
# Normalization using training information
test_data = test_data - train_mean
test_data = test_data / train_std
print "Loading finished..."
print "Sizes...", train_data.shape, train_label.shape, test_data.shape, test_label.shape
return train_data, train_label, test_data, test_label

def load_car_dataset():
    '''
        Load the car dataset in the following format:
        - each image is converted to gray scale
        - each car image is vectorized as a column vector
        - labels are organized as a column vector
    '''
    print "Loading car dataset..."
    # Process training images first
    train_path = './car-dataset/train/'
    train_pos_files = [f for f in os.listdir(os.path.join(train_path, 'positive'))]
    train_neg_files = [f for f in os.listdir(os.path.join(train_path, 'negative'))]
    num_train = len(train_pos_files + train_neg_files)
    train_pos_data = zeros((40*20, len(train_pos_files)))
    train_neg_data = zeros((40*20, len(train_neg_files)))
    for i,f in enumerate(train_pos_files):
        image = imread(os.path.join(train_path, 'positive', f))
        image = cv2.cvtColor(image, cv2.COLOR_RGBA2GRAY)
        image = get_integral_image(image)
        train_pos_data[:,i] = image.flatten()
    for i,f in enumerate(train_neg_files):
        image = imread(os.path.join(train_path, 'negative', f))
        image = cv2.cvtColor(image, cv2.COLOR_RGBA2GRAY)
        image = get_integral_image(image)
        train_neg_data[:,i] = image.flatten()
    # Process testing images

```

```

test_path = './car-dataset/test/'
test_pos_files = [f for f in os.listdir(os.path.join(test_path, 'positive'))]
test_neg_files = [f for f in os.listdir(os.path.join(test_path, 'negative'))]
num_test = len(test_pos_files + test_neg_files)
test_pos_data = zeros((40*20, len(test_pos_files)))
test_neg_data = zeros((40*20, len(test_neg_files)))
for i,f in enumerate(test_pos_files):
    image = imread(os.path.join(test_path, 'positive', f))
    image = cv2.cvtColor(image, cv2.COLOR_RGBA2GRAY)
    image = get_integral_image(image)
    test_pos_data[:,i] = image.flatten()
for i,f in enumerate(test_neg_files):
    image = imread(os.path.join(test_path, 'negative', f))
    image = cv2.cvtColor(image, cv2.COLOR_RGBA2GRAY)
    image = get_integral_image(image)
    test_neg_data[:,i] = image.flatten()
train_data = hstack((train_pos_data, train_neg_data))
train_label = hstack((ones(train_pos_data.shape[1]), zeros(train_neg_data.shape[1])))
test_data = hstack((test_pos_data, test_neg_data))
test_label = hstack((ones(test_pos_data.shape[1]), zeros(test_neg_data.shape[1])))
print "Loading finished..."
print "Sizes...", train_data.shape, train_label.shape, test_data.shape, test_label.shape
print "Type...", train_data.dtype, train_label.dtype
return train_data, train_label, test_data, test_label

```

```

def PCAvsLDA(max_K):
    train_data, train_label, test_data, test_label = load_face_dataset()
    pca_accu = zeros(max_K)
    lda_accu = zeros(max_K)
    for k in range(max_K):
        pca = PCAClassifier()
        pca.train(train_data, train_label, k+1)
        pca_accu[k] = pca.test(test_data, test_label)
        lda = LDAClassifier()
        lda.train(train_data, train_label, k+1)
        lda_accu[k] = lda.test(test_data, test_label)
    line1, = plot(linspace(1,max_K,num=max_K), pca_accu, '-ro', label='PCA')
    line2, = plot(linspace(1,max_K,num=max_K), lda_accu, '-go', label='LDA')
    legend(handles=[line1, line2], loc=4)
    xlabel('K')
    ylabel('Accuracy')
    show()

```

```

def plot_adaboost():
    F_train = array([0.3049, 0.2088, 0.1456, 0.1303, 0.1200, 0.1041, 0.1018, 0.0910,

```

```

D_train = array([1.0000, 1.0000, 1.0000, 1.0000, 1.0000, 1.0000, 1.0000, 0.9986,
A_train = array([0.7828, 0.8513, 0.8963, 0.9072, 0.9145, 0.9259, 0.9275, 0.9348,
F_test = array([0.3341, 0.2477, 0.1773, 0.1568, 0.1432, 0.1409, 0.1386, 0.1318,
D_test = array([0.9831, 0.9663, 0.9551, 0.9438, 0.9438, 0.9382, 0.9382, 0.9382,
A_test = array([0.7573, 0.8139, 0.8608, 0.8722, 0.8819, 0.8819, 0.8835, 0.8883,
N = len(F_train)
for train, test, name, loc in [(F_train, F_test, 'False Positive Rate', 1),
                                (D_train, D_test, 'Detection Rate', 1),
                                (A_train, A_test, 'Accuracy', 1)]:

    figure()
    line1, = plot(range(1,N+1), train, '-ro', label='Train')
    line2, = plot(range(1,N+1), test, '-go', label='Test')
    legend(handles=[line1, line2], loc=loc)
    xlabel('N')
    ylabel(name + ' (%)')

show()

def main():
    algorithms = ['AdaBoost'] # 'PCA', 'LDA', 'AdaBoost'
    # PCA
    if 'PCA' in algorithms:
        train_data, train_label, test_data, test_label = load_face_dataset()
        pca = PCAClassifier()
        pca.train(train_data, train_label, 60)
        pca.test(test_data, test_label)

    # LDA
    elif 'LDA' in algorithms:
        train_data, train_label, test_data, test_label = load_face_dataset()
        lda = LDAClassifier()
        lda.train(train_data, train_label, 30)
        lda.test(test_data, test_label)

    elif 'AdaBoost' in algorithms:
        num_stages = 10
        num_feats = 20
        train_data, train_label, test_data, test_label = load_car_dataset()
        violajones = CascadedAdaBoostClassifier()
        violajones.set_testing_data(test_data, test_label)
        violajones.train(train_data, train_label, num_stages, num_feats)

if __name__ == '__main__':
    main()

```